



COMP1521

WEEK 9

Announcements



- Lab start this week!
 - Lab9 has been released, due on **Week 10 Monday 12:00:00 (midday)**
- weekly quiz 8 be due **Week 9 Thursday 21:00:00**
- **Assignment2** will due on **Week 10 Friday 18:00:00**
- **Final Exam (25/08/2025)**
- **There will be 2 sessions(Morning/afternoon), only need to attend one of these sessions.**
 - Details.....
 - **Week10's lab!!!!**

Contents

Environments

Recurring Directories

Stat.h

Unicode

Common Environment Variables

| Variable | Description |
|----------|--|
| PATH | Specifies directories to search for executable programs (e.g., ls, gcc). |
| HOME | The current user's home directory path (e.g., /home/yourname). |
| PWD | The present working directory (used when running pwd). |
| USER | The name of the current user. |
| SHELL | The path to the current shell program (e.g., /bin/bash). |
| LANG | Default language and locale setting (e.g., en_US.UTF-8). |
| TERM | Specifies the terminal type (e.g., xterm-256color). |

getenv

- `char *getenv(const char *name);`
- - Retrieves the value of an environment variable.
- - Returns the value as a string, or NULL if not found.
- Example:
- `char *path = getenv("PATH");`

opendir

- `DIR *opendir(const char *name);`
- - Opens a directory stream.
- - Returns a `DIR*` on success, `NULL` on failure.
- Example:
- `DIR *dir = opendir("./");`

readdir

- `struct dirent *readdir(DIR *dirp);`
- - Reads the next directory entry.
- - Returns `struct dirent*` or `NULL` when done.
- Example:
- `while ((entry = readdir(dir)) != NULL) {`
- `printf("%s\n", entry->d_name);`
- }

closedir

- `int closedir(DIR *dirp);`
- - Closes an open directory stream.
- - Returns 0 on success, -1 on failure.

write recurse-sample.c:

a file that recurses through a directory and prints out permissions and filenames of all files with "hello" in its name.

```
#include <stdlib.h>
```

```
char *getenv(const char *name);
```

- search environment variable array for **name=value**
- returns **value**
- returns **NULL** if **name** not in environment variable array

```
int main(void) {  
    // print value of environment variable STATUS  
    char *value = getenv("STATUS");  
    printf("Environment variable 'STATUS' has value '%s'\n", value);  
}
```

[source code for get_status.c](#)

write recurse-sample.c:

a file that recurses through a directory and prints out permissions and filenames of all files with "hello" in its name.

```
// part 0: get and print out current directory  
pathname.
```

write recurse-sample.c:

a file that recurses through a directory and prints out permissions and filenames of all files with "hello" in its name.

```
// part 0: get and print out current directory  
pathname.
```

```
// part 1: print out all files in current directory.
```

Searching Directories - Readdir

```
22      // opens DIR pointer
23      DIR *dirp = opendir(current_path);
24      if (dirp == NULL) {
25          perror(current_path);
26          exit(1);
27      }
28
29      struct dirent *de;
30      //loops through all files in DIR, populating struct direct de.
31      while ((de = readdir(dirp)) != NULL) {
32          printf("%s\n", de->d_name);
33      }
```

write recurse-sample.c:

a file that recurses through a directory and prints out permissions and filenames of all files with "hello" in its name.

// part 0: get and print out current directory pathname.

// part 1: print out all files in current directory.

*// part 2: recurse through all subdirectories and print out
pathnames.*

Recurring through Directories -

- We want to Search through each subdirectory using readdir().

Steps:

- Search through current directory using readdir()
- Skip current (.) and parent directories (..)
- If readdir() returns a directory, recursively search through it.
 - use stat.h to find S_IFDIR.
 - Append filename to old pathname to create next directory's pathname.

```
while ((de = readdir(dirp)) != NULL) {  
    printf("%s\n", de->d_name);  
  
    //check if is self or parent  
    char *self = ".";  
    char *parent = "..";  
    if (!strcmp(de->d_name, self) || !strcmp(de->d_name, parent)) {  
        continue;  
    }  
    //if directory, recurse into it.  
    int new_length = strlen(de->d_name) + strlen(current_path) + 2;  
    char *new_path = malloc(new_length * sizeof(char));  
    sprintf(new_path, "%s/%s", current_path, de->d_name);  
    struct stat s;  
    // use lstat  
    lstat(new_path, &s);  
    mode_t perms = s.st_mode;  
    if((perms & S_IFMT) == S_IFDIR) {  
        recurse(new_path);  
    }  
}
```

Stat.h

quick reference guide :

[stat.h quick reference](#)

man 2 stat = stat(pathname, &s) reference

man 7 inode = permissions

reference

Usage:

```
struct stat s;
if (stat(pathname, &s) != 0) {
    perror(pathname);
    exit(1);
}

printf("ino = %10ld # Inode number\n", s.st_ino);
printf("mode = %10o # File mode \n", s.st_mode);
printf("nlink =%10ld # Link count \n", (long)s.st_nlink);
printf("uid = %10u # Owner uid\n", s.st_uid);
printf("gid = %10u # Group gid\n", s.st_gid);
printf("size = %10ld # File size (bytes)\n", (long)s.st_size);

printf("mtime =%10ld # Modification time (seconds since 1/1/70)\n",
       (long)s.st_mtime);
```


stat.st_mode

- Probably the most important part of stat.h alongside length (for 1521)
- Usage:

```
if (s.st_mode & S_IRUSR) {printf("user has read permission");}  
if ((perms & S_IFMT) == S_IFDIR) {printf("This file is a directory!!");}
```

- use this to check for permissions and check if file is a directory!

write recurse-sample.c:
a file that recurses through a directory and prints out permissions and
filenames of all files with "hello" in its name.

// part 0: get and print out current directory pathname.

// part 1: print out all files in current directory.

*// part 2: recurse through all subdirectories and print out
pathnames.*

// part 3: only print out pathnames with "hello" in the filename.

Search for a string

Syntax

Following is the syntax of the C library `strstr()` function –

```
char *strstr (const char *str_1, const char *str_2);
```

Parameters

This function takes two parameters –

- **str_1** – This is a main string.
- **str_2** – The substring to be searched in main string i.e. str_1.

Return Value

The function return the value based on following conditions –

- The function returns a pointer to the first characters of str_2 in str_1 or null pointer if str_2 is not found in str_1.
- If str_2 is found as an empty string, str_1 is returned.

write recurse-sample.c:
a file that recurses through a directory and prints out permissions and
filenames of all files with "hello" in its name.

// part 0: get and print out current directory pathname.

// part 1: print out all files in current directory.

*// part 2: recurse through all subdirectories and print out
pathnames.*

// part 3: only print out pathnames with "hello" in the filename.

// part 4: print out the permissions of these files.

Permissions - s.st_mode

Mode_t (octal):
755

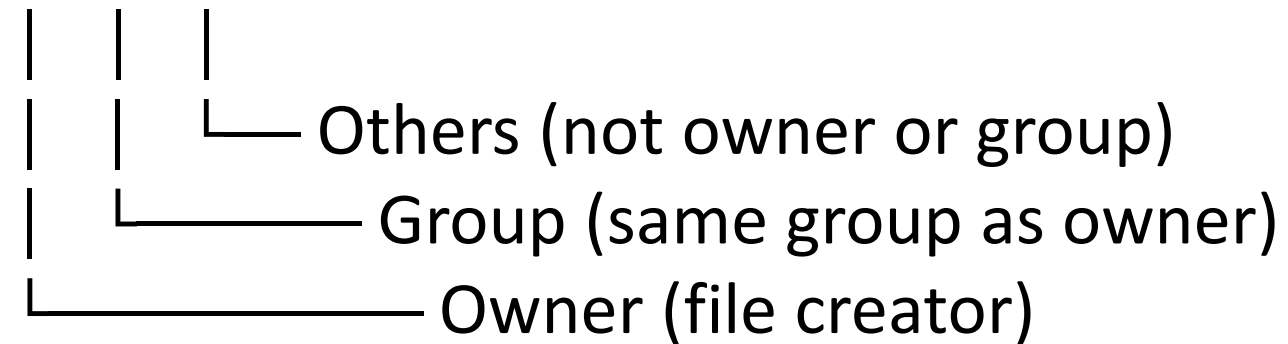
```
if (s.st_mode & S_IRUSR) {printf("user has read permission");}  
if ((perms & S_IFMT) == S_IFDIR) {printf("This file is a directory!!");}
```

```
int main() {  
    struct stat s;  
    stat("a_a", &s);  
    printf("whole mode: %o\n", s.st_mode);  
  
    // only look at last 9  
    uint32_t mask = 0777;  
    uint32_t perms = s.st_mode;  
    perms = perms & mask;  
  
    printf("%o\n", perms);  
    // printf("%b\n", perms);  
    return 0;  
}
```

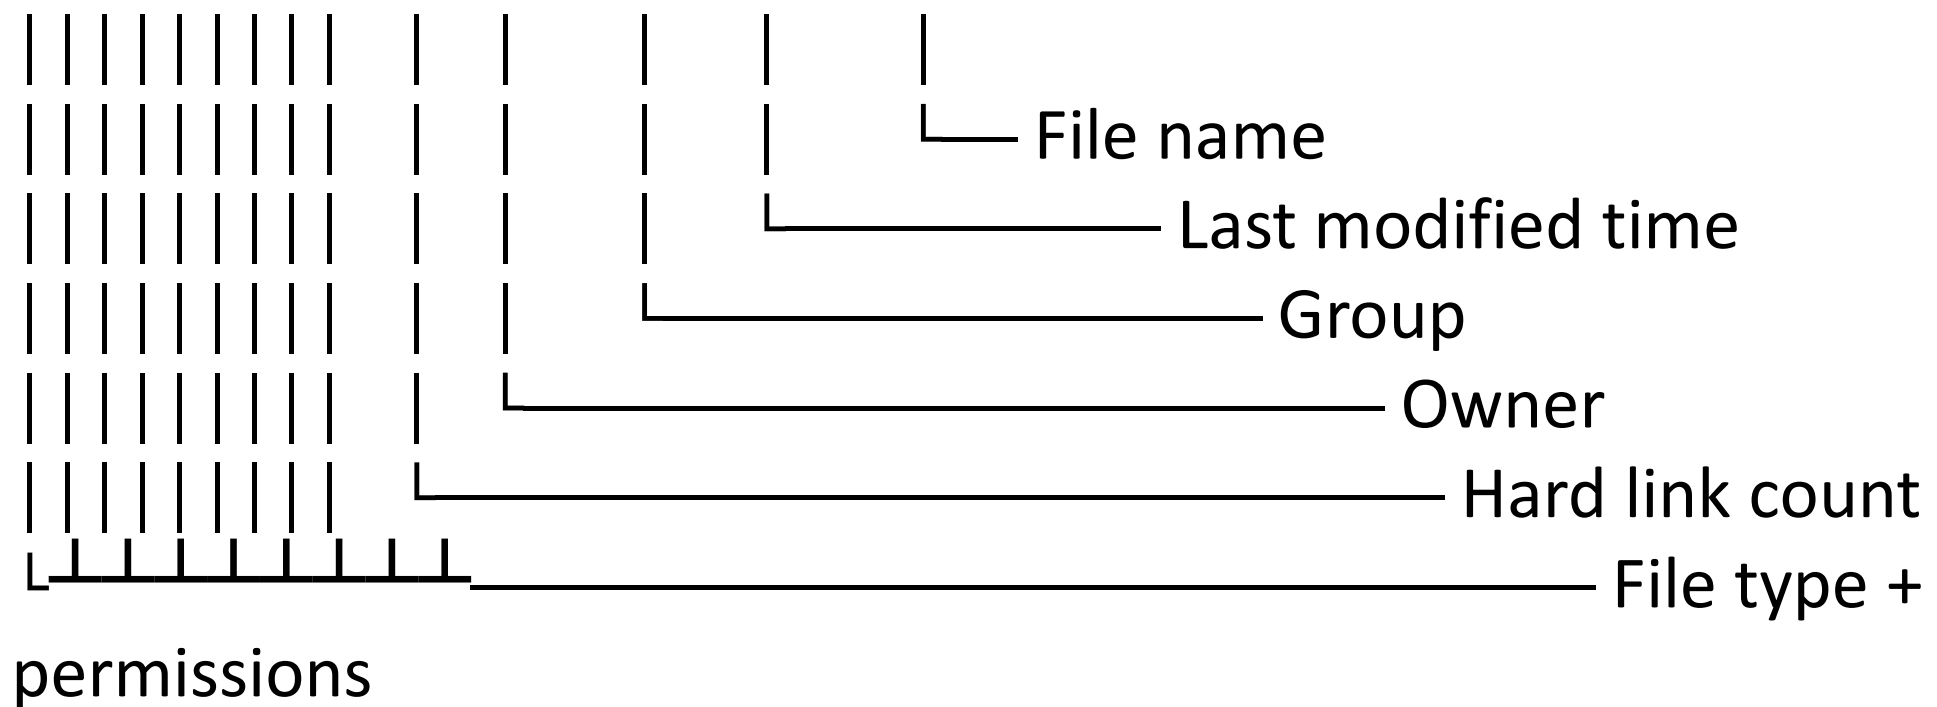
```
z5360323@vx14:~/COMP1521-25T1/tutes/week9$ ./demo  
whole mode: 40755  
755
```

Linux ls -l Permission Format:

rwX r-X r--



-rwxr-xr-- 1 user group 1234 Jul 24 filename



chmod:

What is the meaning for chmod 777?

chmod:

What is the meaning for chmod 777?

| Octal | Binary | Symbolic | Meaning |
|-------|--------|----------|----------------------|
| 7 | 111 | rwX | Read, Write, Execute |
| 6 | 110 | rw- | Read, Write |
| 5 | 101 | r-X | Read, Execute |
| 4 | 100 | r-- | Read only |
| 0 | 000 | --- | No permission |

chmod:

What is the meaning for chmod 777?

| Octal | Binary | Symbolic | Meaning |
|-------|--------|----------|----------------------|
| 7 | 111 | rwX | Read, Write, Execute |
| 6 | 110 | rw- | Read, Write |
| 5 | 101 | r-X | Read, Execute |
| 4 | 100 | r-- | Read only |
| 0 | 000 | --- | No permission |

chmod 777 = Everyone can do anything (rwX)

Linux Permission Macro Matrix

| | Owner | Group | Others |
|------------------|--------------|--------------|---------------|
| Read (r) | S_IRUSR | S_IRGRP | S_IROTH |
| Write (w) | S_IWUSR | S_IWGRP | S_IWOTH |
| Exec (x) | S_IXUSR | S_IXGRP | S_IXOTH |

Permissions - s.st_mode

Mode_t (octal): 755

Binary: 111 101
101

```
if (s.st_mode & S_IRUSR) {printf("user has read permission");}  
if ((perms & S_IFMT) == S_IFDIR) {printf("This file is a directory!!");}
```

```
int main() {  
    struct stat s;  
    stat("sample", &s);  
    printf("whole mode: %o\n", s.st_mode);  
  
    // only look at last 9  
    uint32_t mask = 0777;  
    uint32_t perms = s.st_mode;  
    perms = perms & mask;  
  
    printf("%o\n", perms);  
    printf("%b\n", perms);  
    return 0;  
}
```

```
z5360323@vx14:~/COMP1521-25T1/tutes/week9$ ./demo  
whole mode: 40755  
755  
111101101
```

Permissions - s.st_mode

Mode_t (octal): 755

Binary: 111 101 101

String: rwx r-x r-x

```
if (s.st_mode & S_IRUSR) {printf("user has read permission");}  
if ((perms & S_IFMT) == S_IFDIR) {printf("This file is a directory!!");}
```

```
int main() {  
    struct stat s;  
    stat("a_a", &s);  
    printf("whole mode: %o\n", s.st_mode);  
  
    // only look at last 9  
    uint32_t mask = 0777;  
    uint32_t perms = s.st_mode;  
    perms = perms & mask;  
  
    printf("%o\n", perms);  
    printf("%b\n", perms);  
    return 0;  
}
```

```
z5360323@vx14:~/COMP1521-25T1/tutes/week9$ ./demo  
whole mode: 40755  
755  
111101101  
z5360323@vx14:~/COMP1521-25T1/tutes/week9$ ls -l  
total 216  
drwxr-xr-x 2 z5360323 z5360323 4096 Apr 14 10:05 a_a
```

Permissions - s.st_mode

Mode_t (octal): 755

Binary: 111 101 101

String: rwx r-w r-w

To convert from binary to string,

- print “r/w/x” if bin = 1
- else print “-”.
- each binary digit has a hash define!

```
if (s.st_mode & S_IRUSR) {printf("user has read permission");}  
if ((perms & S_IFMT) == S_IFDIR) {printf("This file is a directory!!");}
```

```
int main() {  
    struct stat s;  
    stat("a_a", &s);  
    printf("whole mode: %o\n", s.st_mode);  
  
    // only look at last 9  
    uint32_t mask = 0777;  
    uint32_t perms = s.st_mode;  
    perms = perms & mask;  
  
    printf("%o\n", perms);  
    printf("%b\n", perms);  
    return 0;  
}
```

```
z5360323@vx14:~/COMP1521-25T1/tutes/week9$ ./demo  
whole mode: 40755  
755  
111101101  
z5360323@vx14:~/COMP1521-25T1/tutes/week9$ ls -l  
total 216  
drwxr-xr-x 2 z5360323 z5360323 4096 Apr 14 10:05 a_a
```

Permissions - s.st_mode

```
if (s.st_mode & S_IRUSR) {printf("user has read permission");}  
if ((perms & S_IFMT) == S_IFDIR) {printf("This file is a directory!!");}
```

| | |
|----|--|
| 55 | (perms & S_IRUSR) ? printf("r") : printf("-"); |
| 56 | (perms & S_IWUSR) ? printf("w") : printf("-"); |
| 57 | (perms & S_IXUSR) ? printf("x") : printf("-"); |
| 58 | (perms & S_IRGRP) ? printf("r") : printf("-"); |
| 59 | (perms & S_IWGRP) ? printf("w") : printf("-"); |
| 60 | (perms & S_IXGRP) ? printf("x") : printf("-"); |
| 61 | (perms & S_IROTH) ? printf("r") : printf("-"); |
| 62 | (perms & S_IWOTH) ? printf("w") : printf("-"); |
| 63 | (perms & S_IXOTH) ? printf("x") : printf("-"); |

64

What is UTF-8?

- UTF-8 stands for 8-bit Unicode Transformation Format
- It's a variable-length encoding for Unicode codepoints
- Each character (codepoint) is encoded using 1 to 4 bytes
- Compatible with ASCII (0x00 to 0x7F)

UTF-8 (cont.)

This is the layout of UTF-8

| #bytes | #bits | Byte 1 | Byte 2 | Byte 3 | Byte 4 |
|--------|-------|----------|----------|----------|----------|
| 1 | 7 | 0xxxxxxx | - | - | - |
| 2 | 11 | 110xxxxx | 10xxxxxx | - | - |
| 3 | 16 | 1110xxxx | 10xxxxxx | 10xxxxxx | - |
| 4 | 21 | 11110xxx | 10xxxxxx | 10xxxxxx | 10xxxxxx |

A single UTF-8 character can be anywhere from 1 to 4 bytes long

The entire ASCII character set can be represented in 1 byte with zero wasted bits

The entire BMP can be represented in 3 bytes, being 8 bits more efficient than UTF-32

and the entire UNICODE character set can be represented in 4 bytes/ using exactly the same number of bits as UTF-32 in the worst case

Bitmask Explanation

| #bytes | #bits | Byte 1 | Byte 2 | Byte 3 | Byte 4 |
|--------|-------|----------|----------|----------|----------|
| 1 | 7 | 0xxxxxxx | - | - | - |
| 2 | 11 | 110xxxxx | 10xxxxxx | - | - |
| 3 | 16 | 1110xxxx | 10xxxxxx | 10xxxxxx | - |
| 4 | 21 | 11110xxx | 10xxxxxx | 10xxxxxx | 10xxxxxx |

| Mask | Matches Pattern | Meaning |
|------|-----------------|----------------|
| 0x80 | 0xxxxxxx | 1-byte (ASCII) |
| 0xE0 | 110xxxxx | 2-byte start |
| 0xF0 | 1110xxxx | 3-byte start |
| 0xF8 | 11110xxx | 4-byte start |

How to Determine Character Length in C

```
int utf8_char_length(unsigned char byte) {  
    if ((byte & 0x80) == 0) return 1;           // 0x80 = 10000000  
                                                // 0xC0 = 11000000  
  
    else if ((byte & 0xE0) == 0xC0) return 2;    // 0xE0 = 11100000  
    else if ((byte & 0xF0) == 0xE0) return 3;    // 0xF0 = 11110000  
    else if ((byte & 0xF8) == 0xF0) return 4;    // 0xF8 = 11111000  
    else return -1;  
}
```

Tut.q8

View the number of characters (codepoints)

```
echo -n "早上好中国现在我有冰淇淋" | wc -m
```

View the number of bytes (total length after UTF-8 encoding):

```
echo -n "早上好中国现在我有冰淇淋" | wc -c
```

ANSWER:

The `struct stat` fields are:

`ino_t st_ino`

An inode number, giving an index into the filesystem's table of file metadata structures. For `stat.c`, it could be any largish positive integer. The inode number can be accessed using `ls -li`.

`mode_t st_mode`

Contains information about the file type and the file permissions, encoded as a bit-string. These bit-strings are usually written in octal, to make it easy to see the 3 groups of 3 bits defining the file permissions. A regular file like `stat.c` would have an `st_mode` value of `0100644` (from `S_IFREG | S_IRUSR | S_IWUSR | S_IRGRP | S_IROTH`).

`uid_t st_uid`

Gives the numeric user id (uid) of the user to whom the file belongs (its owner); in this case, 516. This can be retrieved using `ls -ln`.

`gid_t st_gid`

Gives the numeric id (gid) of the group to which the file belongs; in this case, 36820. This can be retrieved using `ls -ln`.

`off_t st_size`

Gives the total size of the file in bytes. For a text file like `stat.c`, it's simply the number of characters in the file's content (i.e., 1855).

`blksize_t st_blksize`

Gives the size of a block on the storage device useful for files of this type. Typical block size are 512, 1024, 4096, 8192.

`blkcnt_t st_blocks`

Gives the amount of space on the storage device allocated for this file. Since it's allocated in 512B chunks, more space might be allocated than is actually required to store the bytes. Often blocks are allocated in groups of size 2^n . The total bytes allocated in the blocks must, of course, be larger than `st_size`. For `stat.c`, there are 8 blocks allocated (a total of 4096 bytes, to store the 1855 actually in the file).

`time_t st_mtime`

Gives the last time the file was modified. A `time_t` value is typically implemented as an integer giving the number of seconds since midnight on Jan 1 1970. For the `stat.c` file, the most recent update time is shown in the `ls` output as `Sep 9 14:24`, which, here, is implied to be `2017/09/09 14:24`, a value around 1504931040.

`time_t st_atime`

The last time the file content was accessed (read or written). This value can be retrieved using `ls -lu`.

`time_t st_ctime`

The last time the file status was changed. This could mean changing the file contents, or changing its associated metadata. This value can be retrieved using `ls -lc`.